

Databases: Supervision 1

Daniel Chatfield

February 15, 2015

2003 Paper 5 Question 8

1. (a) i. Define the operators in the core relational algebra. [5]

Union $R \cup S$ is the union of R and S .

Intersection $R \cap S$ is the intersection of R and S .

Difference $R - S$ is the difference between R and S .

Product $R \times S$ is the product of R and S .

Selection $\sigma_p(R)$ is a selection over R . P is a propositional formula and the operator selects all tuples in R for which p holds.

Projection $\pi_x(R)$ is a projection over R . x is a set of attribute names.

- ii. Define the domain relational calculus. [4]

$$Q = \{(A_1 = v_1, A_2 = v_2, \dots, A_k = v_k) | P(v_1, v_2, \dots, v_k)\}$$

- iii. Show how the relational algebra can be encoded in the domain relational calculus. [3]

- Selection

Relational Algebra $Q = \sigma_{A > 12}(R)$

Domain Relational Calculus Encoded as:

$$Q = \{ \{ (A, a), (B, b), (C, c) \} | \{ (A, a), (B, b), (C, c) \} \in R \wedge a > 12 \}$$

- Projection

Relational Algebra $Q = \pi_{B,C}(R)$

Domain Relational Calculus Encoded as:

$$Q = \{ \{ (B, b), (C, c) \} \mid \{ (A, a), (B, b), (C, c) \} \in R \}$$

- (b) A *constraint* can be expressed using relational algebra. For example, $R = \emptyset$ specifies the constraint that relation R must be empty, and $(R \cup S) \subseteq T$ specifies that every tuple in the union of R and S must be in T .

Consider the following schema.

RockStar (name, address, gender, birthday)
RockManager (managername, starname)

- i. Give a constraint to express that rock stars must be either male or female. [1]

$$\sigma_{gender \neq 'male' \text{ AND } gender \neq 'female'}(RockManager) = \emptyset$$

- ii. Give a constraint to express the referential integrity constraint between the RockStar and RockManager relations. (Note: starname is intended to be a foreign key.) [3]

$$\pi_{starname}(RockManager) \subseteq \pi_{name}(RockStar)$$

- iii. Give a constraint to express the functional dependency $name \rightarrow address$ for the RockStar relation. [4]

Not covered this yet and I don't understand it from slides alone

2004 Paper 5 Question 8

Assume a simple movie database with the following schema. (You may assume that producers have a unique certification number, *Cert*, that is also recorded in the *Movie* relation as attribute *prodC#*; and no two movies are produced with the same title.)

Movie (title, year, length, prodC#)
StarsIn (movieTitle, movieYear, starName)
Producer (name, address, cert)
MovieStar (name, gender, birthdate)

2. (a) Write the following queries in SQL:

i. Who were the male stars in the film *The Red Squirrel*? [1]

```
(SELECT name
FROM MovieStar
WHERE gender='male')
INTERSECT
(SELECT starName
FROM StarsIn
WHERE movieTitle='The Red Squirrel')
```

ii. Which movies are longer than *Titanic*? [2]

```
SELECT *
FROM Movie
WHERE length >
  (SELECT length
   from Movie
   WHERE title='Titanic')
```

(b) SQL has a boolean-valued operator `IN` such that the expression `s IN R` is true when `s` is contained in the relation `R` (assume for simplicity that `R` is a single attribute relation and hence `s` is a simple atomic value).

Consider the following nested SQL query that uses the `IN` operator:

```
SELECT name
FROM Producer
WHERE cert IN (SELECT prodC#
               FROM Movie
               WHERE title IN (SELECT movieTitle
                              FROM StarsIn
                              WHERE starName='Nancho Novo'));
```

i. State concisely what this query is intended to mean. [1]

This query returns the names of the producers of the movies that 'Nancho Novo' starred in.

ii. Express this nested query as a single `SELECT-FROM-WHERE` query. [2]

```
SELECT DISTINCT Producer.name
FROM StarsIn, Movie, Producer
WHERE StarsIn.starName = 'nancho Novo',
      Movie.title = StarsIn.movieTitle,
      Producer.cert = Movie.prodC#
```

- iii. Is your query from part (b)(ii) always equivalent to the original query? If yes, then justify your answer; if not, then explain the difference and show how they could be made equivalent. [6]

The query is always equivalent, the `DISTINCT` ensures that a producer is only returned once, just like the original query.

The outer-most part of the original query was over the Producer table and therefore a single producer could only be returned once, even if they produced multiple films with Nancho Novo since the predicate `IN` is simply *true* or *false*.

The new query joins the tables together which duplicates the Producer field and therefore the `DISTINCT` is required to ensure no duplicates.

- (c) SQL has a boolean-valued operator `EXISTS` such that `EXISTS R` is true if and only if `R` is not empty. [8]

Show how `EXISTS` is, in fact, redundant by giving a simple SQL expression that is equivalent to `EXISTS R` but does not involve `EXISTS` or any cardinality operators, e.g. `COUNT`. [Hint: You may use the `IN` operator.]

I think this is a valid solution but the number of marks available for this question is making me doubt that.

`R IN R`

An example:

```
SELECT *
FROM Colleges
WHERE EXISTS (
  SELECT *
  FROM Students
  WHERE Colleges.id = Students.college
)
```

can be written as:

```
SELECT *
FROM Colleges
WHERE id IN (
  SELECT college
  FROM Students
)
```